




Optimal control for **walking** robots


Theory and **practice** with Crocoddyl

Nicolas Mansard

Gepetto, LAAS-CNRS & ANITI



#04: Differential Dynamic Programming




Multiple views on DDP



1. DDP as iterative LQR


DDP as iterative LQR



- “Next-step” is a nonlinear function

$$\Delta x' = f(x + \Delta x, u + \Delta u) - f(x, u)$$

DDP as iterative LQR



- “Next-step” is a nonlinear function

$$\Delta x' = f(x + \Delta x, u + \Delta u) - f(x, u)$$

- Approximate by

$$\Delta x' = f(x, u) + F_x \Delta x + F_u \Delta u - f(x, u)$$

DDP as iterative LQR

- Nonlinear optimal control problem

$$\begin{aligned} \min_{\{x\}, \{u\}} \quad & \sum_{t=0}^{T-1} l(x_t, u_t) + l_T(x_T) \\ \text{s.t.} \quad & \forall t=0..T-1 \quad x_{t+1} = f(x_t, u_t) \end{aligned}$$

- Linear-Quadratic problem ... solved with Ricatti recursion (textbook)

$$\begin{aligned} \min_{\{\Delta x\}, \{\Delta u\}} \quad & \sum_{t=0}^{T-1} \begin{pmatrix} L_x \\ L_u \end{pmatrix}^T \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \frac{1}{2} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix}^T \begin{pmatrix} L_{xx} & L_{xu} \\ L_{ux} & L_{uu} \end{pmatrix} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \dots \\ \text{s.t.} \quad & \forall t=0..T-1 \quad \Delta x_{t+1} = F_x \Delta x_t + F_u \Delta u_t \end{aligned}$$

DDP as iterative LQR

- Algorithm iLQR

Initialize with a given trajectory $\{\mathbf{x}_0\}, \{\mathbf{u}_0\}$

Repeat

 Linearize/Quadratize the OCP

 Compute the LQR policy

 Simulate (roll-out) with LQR regulator

Until local minimum is reached

Multiple views on DDP



2. DDP as a 2-pass algorithm

DDP as a 2-pass algorithm

$$V_t = \min_{u_t} l(x_t, u_t) + V_{t+1}(f(x_t, u_t))$$

- Backward propagation

$$Q_t = l(x_t, u_t) + V_{t+1}(f(x_t, u_t))$$

- Greedy optimization

$$V_t = \min_{u_t} Q_t(x_t, u_t)$$

DDP as a 2-pass algorithm

$$Q = l + V'$$

$$V = \min_u Q$$

DDP as a 2-pass algorithm

- Pass 1: back-propagate an approximation of V
 - We can solve Bellman for quadratic cost and linear dynamics
- Pass 2: forward propagate gains and trajectory

DDP as a 2-pass algorithm

- Pass 1: backpropagate an approximation of V

DDP as a 2-pass algorithm

- Pass 2: forward propagate gains and trajectory

DDP as a 2-pass algorithm

- Globalization (because nonconvexity)
- Line search
 - $u = u^* + k + K (x - x^*)$
 - $x' = f(x, u)$
- Regularization
 - $Q_{uu} = L_{uu} + F_u^T V_{xx} F_u$
 - $k = Q_{uu}^{-1} Q_u$
 - $K = Q_{uu}^{-1} Q_{ux}$

Multiple views on DDP



3. DDP as sparse SQP

DDP as sparse SQP

$$\begin{aligned} & \min_{\{x\}, \{u\}} \sum_{t=0}^{T-1} l(x_t, u_t) + l_T(x_T) \\ \text{s.t. } & \forall t = 0..T-1 \quad x_{t+1} = f(x_t, u_t) \end{aligned}$$

DDP as sparse SQP

- Reminder
- Non linear problem

$$\begin{aligned} \min_y \quad & l(y) \\ \text{s.t.} \quad & f(y)=0 \end{aligned}$$

- Resulting “linearization”

$$\begin{aligned} \min_{\Delta y} \quad & l(y) + L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ \text{s.t.} \quad & f(y) + F_y \Delta y = 0 \end{aligned}$$

DDP as sparse SQP

$$\begin{aligned} \min_{\Delta y} \quad & l(y) + L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ \text{s.t.} \quad & f(y) + F_y \Delta y = 0 \end{aligned}$$


- Lagrangian on the NLP

$$\mathcal{L}(y, \lambda) = l(y) + \lambda^T f(y)$$

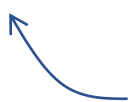
Lagrangian



Primal variable



Dual variable (multipliers)



DDP as sparse SQP

$$\begin{aligned} \min_{\Delta y} \quad & l(y) + L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ \text{s.t.} \quad & f(y) + F_y \Delta y = 0 \end{aligned}$$

- Lagrangian on the QP

$$\begin{aligned} \mathcal{L}(\Delta y, \lambda) = & L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ & + \lambda^T (F_y \Delta y - f(y)) \end{aligned}$$

DDP as sparse SQP

$$\begin{aligned} \min_{\Delta y} \quad & l(y) + L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ \text{s.t.} \quad & f(y) + F_y \Delta y = 0 \end{aligned}$$

- Lagrangian on the QP

$$\begin{aligned} \mathcal{L}(\Delta y, \lambda) = & L_y \Delta y + \frac{1}{2} \Delta y^T L_{yy} \Delta y \\ & + \lambda^T (F_y \Delta y - f(y)) \end{aligned}$$

- Newton step

$$\begin{pmatrix} L_{yy} & F_y^T \\ F_y & 0 \end{pmatrix} \begin{pmatrix} \Delta y \\ \lambda \end{pmatrix} = \begin{pmatrix} -L_y \\ -f(y) \end{pmatrix}$$

DDP as sparse SQP

$$\min_{\{\Delta x\}, \{\Delta u\}} \sum_{t=0}^{T-1} \begin{pmatrix} L_x \\ L_u \end{pmatrix}^T \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \frac{1}{2} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix}^T \begin{pmatrix} L_{xx} & L_{xu} \\ L_{ux} & L_{uu} \end{pmatrix} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \dots$$

$$\text{s.t. } \forall t=0..T-1 \quad \Delta x_{t+1} = F_x \Delta x_t + F_u \Delta u_t + f_t$$

$$\begin{pmatrix} L_{yy} & F_y^T \\ F_y & 0 \end{pmatrix} \begin{pmatrix} \Delta y \\ \lambda \end{pmatrix} = \begin{pmatrix} -L_y \\ -f(y) \end{pmatrix}$$

DDP as sparse SQP

$$\min_{\{\Delta x\}, \{\Delta u\}} \sum_{t=0}^{T-1} \begin{pmatrix} L_x \\ L_u \end{pmatrix}^T \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \frac{1}{2} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix}^T \begin{pmatrix} L_{xx} & L_{xu} \\ L_{ux} & L_{uu} \end{pmatrix} \begin{pmatrix} \Delta x_t \\ \Delta u_t \end{pmatrix} + \dots$$

$$\text{s.t.} \quad \forall t=0..T-1 \quad \Delta x_{t+1} = F_x \Delta x_t + F_u \Delta u_t + f_t$$

$$\begin{bmatrix} L_{xx} & & & L_{xu} & -I & F_x^T \\ & \ddots & & & & \\ & & L_{xx} & & & \\ & & & L_{xu} & & \\ & & & & -I & F_x^T \\ & & & & & -I \\ L_{ux} & & & L_{uu} & F_u^T & \\ & \ddots & & & & \\ & & L_{ux} & & & \\ & & & L_{uu} & & \\ & & & & F_u^T & \\ -I & & & & & \\ F_x & -I & & F_u & & \\ & \ddots & \ddots & & & \\ & & F_x & -I & & \\ & & & & F_u & \end{bmatrix} \begin{bmatrix} \Delta x_0 \\ \vdots \\ \Delta x_{T-1} \\ \Delta x_T \\ \Delta u_0 \\ \vdots \\ \Delta u_{T-1} \\ \lambda_0 \\ \lambda_1 \\ \vdots \\ \lambda_{T-1} \end{bmatrix} = - \begin{bmatrix} L_x \\ \vdots \\ L_x \\ L_x \\ L_u \\ \vdots \\ L_u \\ f_0 \\ f_1 \\ \vdots \\ f_{T-1} \end{bmatrix}$$

Stagewise Implementations of Sequential Quadratic Programming for Model-Predictive Control

Armand Jordana^{*1}, Sébastien Kleff^{*1}, Avadesh Meduri^{*1},
Justin Carpentier², Nicolas Mansard³ and Ludovic Righetti¹

Abstract—The promise of model-predictive control in robotics has led to extensive development of efficient numerical optimal control solvers in line with differential dynamic programming because it exploits the sparsity induced by time. In this work, we argue that this effervescence has hidden the fact that sparsity can be equally exploited by standard nonlinear optimization. In particular, we show how a tailored implementation of sequential quadratic programming achieves state-of-the-art model-predictive control. Then, we clarify the connections between popular algorithms from the robotics community and well-established optimization techniques. Further, the sequential quadratic program formulation naturally encompasses the constrained case, a notoriously difficult problem in the robotics community. Specifically, we show that it only requires a sparsity-exploiting implementation of a state-of-the-art quadratic programming solver. We illustrate the validity of this approach in a comparative study and experiments on a torque-controlled manipulator. To the best of our knowledge, this is the first demonstration of nonlinear model-predictive control with arbitrary constraints on real hardware.

I. INTRODUCTION

A. Motivation

Model Predictive Control (MPC) has become popular for online robot decision-making. It has shown compelling results with all kinds of robots ranging from industrial manipulators [1], quadrupeds [2]–[4] to humanoids [5], [6]. The general idea of MPC is to formulate the robot motion generation problem as a numerical optimization problem, i.e., a finite horizon Optimal Control Problem (OCP), and solve it online at every control cycle using the current state measurement as the initial state. This receding horizon strategy allows us to adapt the robot behavior as the state of the system and environment change.

In robotics, Differential Dynamic Programming (DDP) [7] is a popular choice to solve OCPs because it exploits the problem's structure well. This advantage has led to a bustling algorithmic development over the past two decades [8]–[20]. In light of the increasing number of variations of DDP, one might naively ask: why not use well-established optimization algorithms [21]? Is there anything special in MPC that cannot

be tackled by, for example, an efficient Sequential Quadratic Programming (SQP) solver? In this paper, we show that special implementations of SQP methods developed by the optimization-based community [23]–[26] are, in fact, sufficient to achieve state-of-the-art MPC on real robots.

B. Related work

Mayne first introduced DDP [7] as an efficient method to solve nonlinear OCPs by iteratively performing a backward pass over the time horizon and a forward rollout of the dynamics. This algorithm has a linear complexity in the time horizon and guarantees convergence [27]. More recently, Todorov et al. introduced in DDP by proposing the iterative Linear Quadratic Regulator (iLQR) [8], a variant discarding the second-order terms of the dynamics. It has since gained a lot of traction in the robotics community [3]–[5], [14], [18], [19]. To Gauss-Newton optimization has been applied to DDP [28]. However, this approach faces two limitations: 1) as a single shooting method, it requires a feasible initial guess, which makes the algorithm sensitive to warm-start, an essential requirement for long-horizon problems [12] and 2) enforcing equality constraints is not straightforward. The cost of enforcing constraints softly using penalty functions is not straightforward. But this approach is heuristic (i.e., weight tuning) and tends to cause numerical issues.

Multiple shooting for optimal control, in contrast, addresses the first limitation: it accepts an arbitrary initial guess. Several multiple shooting variants of DDP have been proposed in [12], [14] with significantly improved convergence abilities, which have enabled nonlinear MPC at high frequency on real robots [1], [3], [6], [14].

The second issue of enforcing constraints inside a DDP-like algorithm has been addressed in several works. [10] uses a DDP-based projected Newton method to bound control inputs. This approach has further been improved and deployed on a real quadruped robot in [17]. More recently, augmented Lagrangian methods have been used to enforce constraints in iLQR/DDP algorithms [11], [13], [16], [19]. However, their convergence behavior is less understood than DDP, whose seminal paper [7] was followed by sophisticated proofs [27]. To the best of our knowledge, it has not yet been shown that those recent DDP-based algorithms exhibit global convergence (i.e., convergence from any initial point to a stationary point) and quadratic local convergence.

$$\begin{bmatrix} \Gamma_1 & M_1^T & 0 & 0 & \cdots & 0 \\ M_1 & \Gamma_2 & M_2^T & 0 & \cdots & 0 \\ 0 & M_2 & \Gamma_3 & M_3^T & \cdots & 0 \\ 0 & 0 & M_3 & \Gamma_4 & \ddots & 0 \\ \vdots & \vdots & \vdots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & 0 & \ddots & \Gamma_T \end{bmatrix} \begin{bmatrix} s_1 \\ s_2 \\ s_3 \\ s_4 \\ \vdots \\ s_T \end{bmatrix} = \begin{bmatrix} g_1 \\ g_2 \\ g_3 \\ g_4 \\ \vdots \\ g_T \end{bmatrix} \quad (7)$$

where:

$$\Gamma_k = \begin{bmatrix} R_{k-1} & 0 & -B_{k-1}^T \\ 0 & Q_k & I \\ -B_{k-1} & I & 0 \end{bmatrix}, \quad M_k = \begin{bmatrix} 0 & S_k^T & 0 \\ 0 & 0 & 0 \\ 0 & -A_k & 0 \end{bmatrix}$$

$$\text{and } s_k = \begin{bmatrix} \Delta u_{k-1} \\ \Delta x_k \\ -\lambda_k \end{bmatrix}, \quad g_k = \begin{bmatrix} -r_{k-1} \\ -q_k \\ \gamma_k \end{bmatrix}. \quad (8)$$


This work was in part supported by the National Science Foundation grants 1932187, 2026479, 2222815 and 2315396.

^{*} Equal contribution - first authors listed in alphabetical order.


¹ Machines in Motion Laboratory, New York University, USA. {aj2988@nyu.edu, sk8001@nyu.edu, am9789@nyu.edu, lr114@nyu.edu}

² Inria - Département d'Informatique de l'École normale supérieure, PSL Research University. justin.carpentier@inria.fr

³ LAAS-CNRS, Université de Toulouse, CNRS, Toulouse. nmansard@laas.fr



Sébastien
Kleff



Armand
Jordana



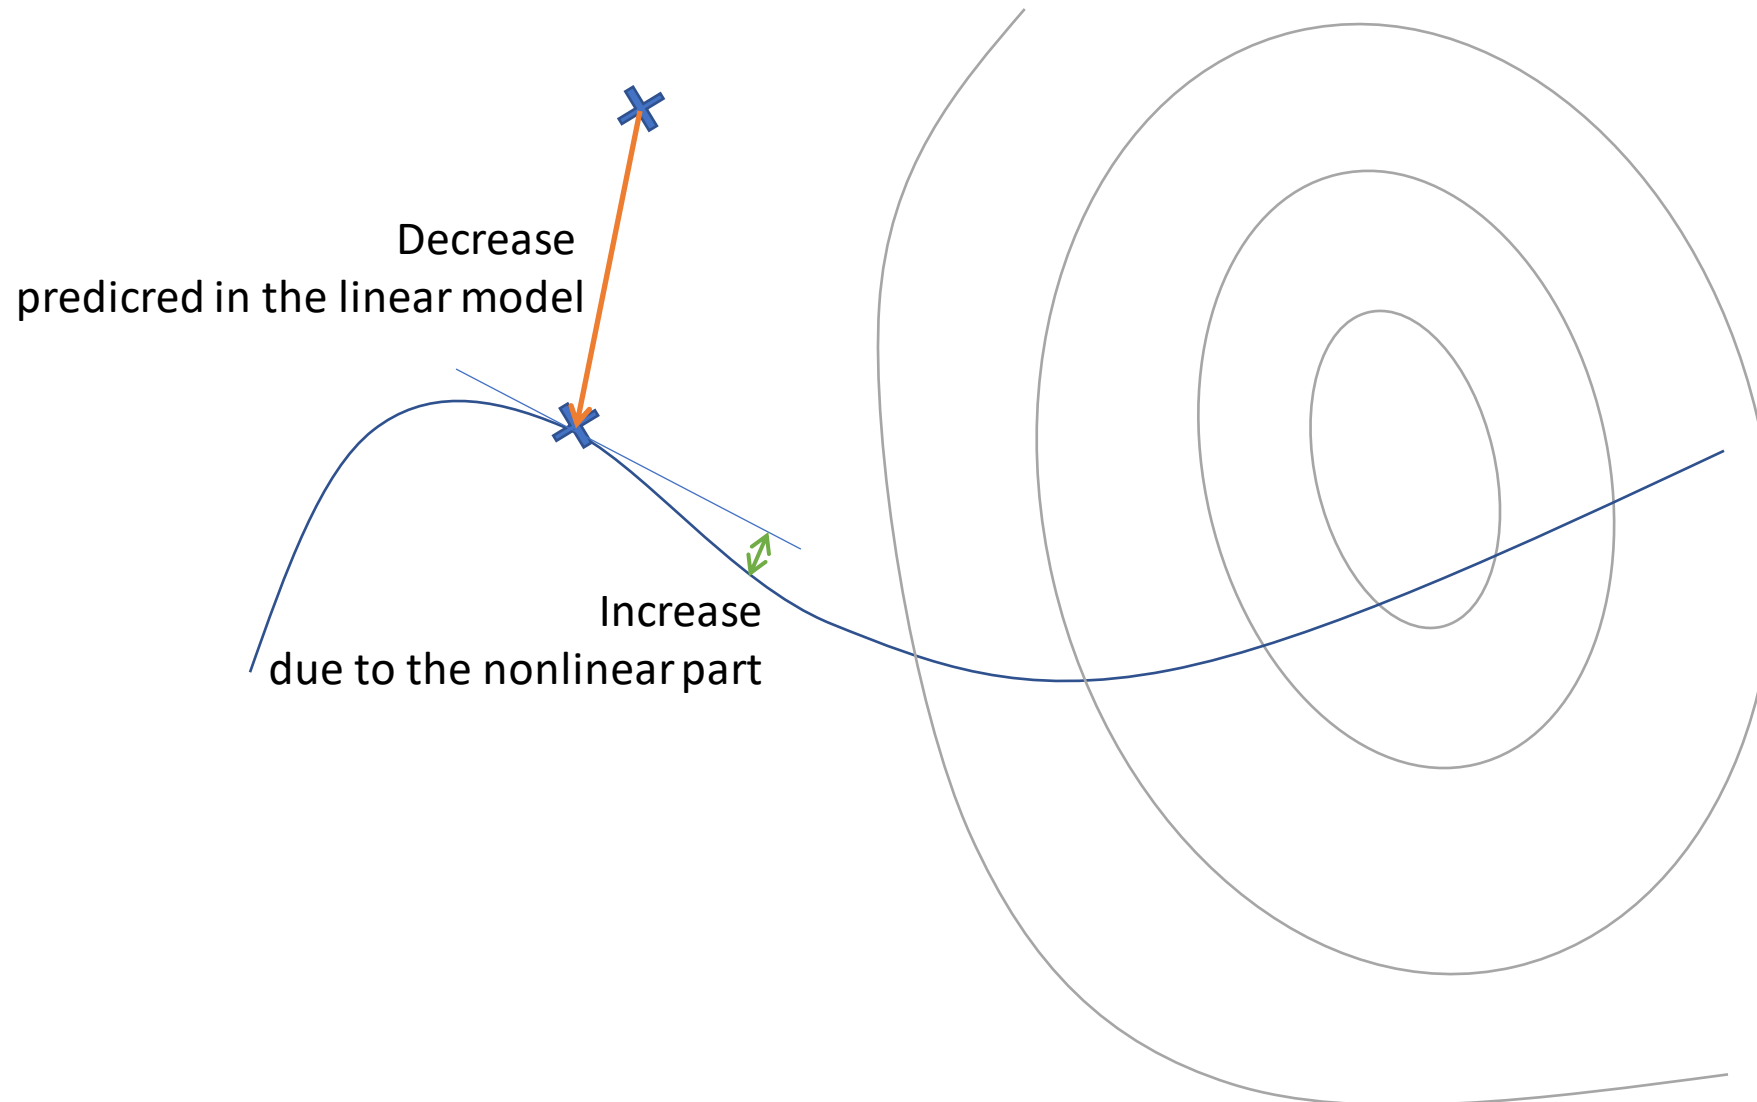
Avadesh
Meduri



Rohan
Budhiraja




Linear versus Nonlinear Rollouts



Linear versus Nonlinear Rollouts


- Case where the current guess is feasible: $x_+ = f(x, u)$
- If we know $\{u\}$, then we can get a feasible X by roll-out

$$x_+ + \Delta x := f(x, u + \alpha \Delta u)$$



Linear versus Nonlinear Rollouts

- Case where the current guess is not feasible: $x_{t+1} = f(x_t, u_t) + f_{t+1}$



Feasibility DDP: a (somehow) multiple-shooting algorithm

- Given the OCP formulated as an NLP
- Iterates:
 - Linearize the NLP ... obtain a cQP
 - Compute a partial step using a LQR
 - Backward pass to backpropagate the Riccati gains
 - Increase regularization if Hamiltonian Hessian becomes nonpositive
 - Perform a line-search in the direction $\Delta X, \Delta U$
 - For any step α , perform a forward propagation with gaps (defect) reduced of α
 - Accept under Wolfe conditions
 - Stop when Lagrangian gradients vanish

Policy derivatives

- At convergence,
the derivatives of the NLP and the LCQP are the same
- The linear optimal policy (LQ regulator) is

$$\Delta\pi(\Delta x) = k + K\Delta x$$

- The derivatives of the nonlinear optimal policy is

$$\frac{\partial\pi}{\partial x} = K$$

Policy Taylor approximation

- A approximation of the optimal policy is

$$\pi(x + \Delta x) = \pi(x) + K \Delta x$$

- The derivatives of the nonlinear optimal policy is

$$\frac{\partial \pi}{\partial x} = K$$